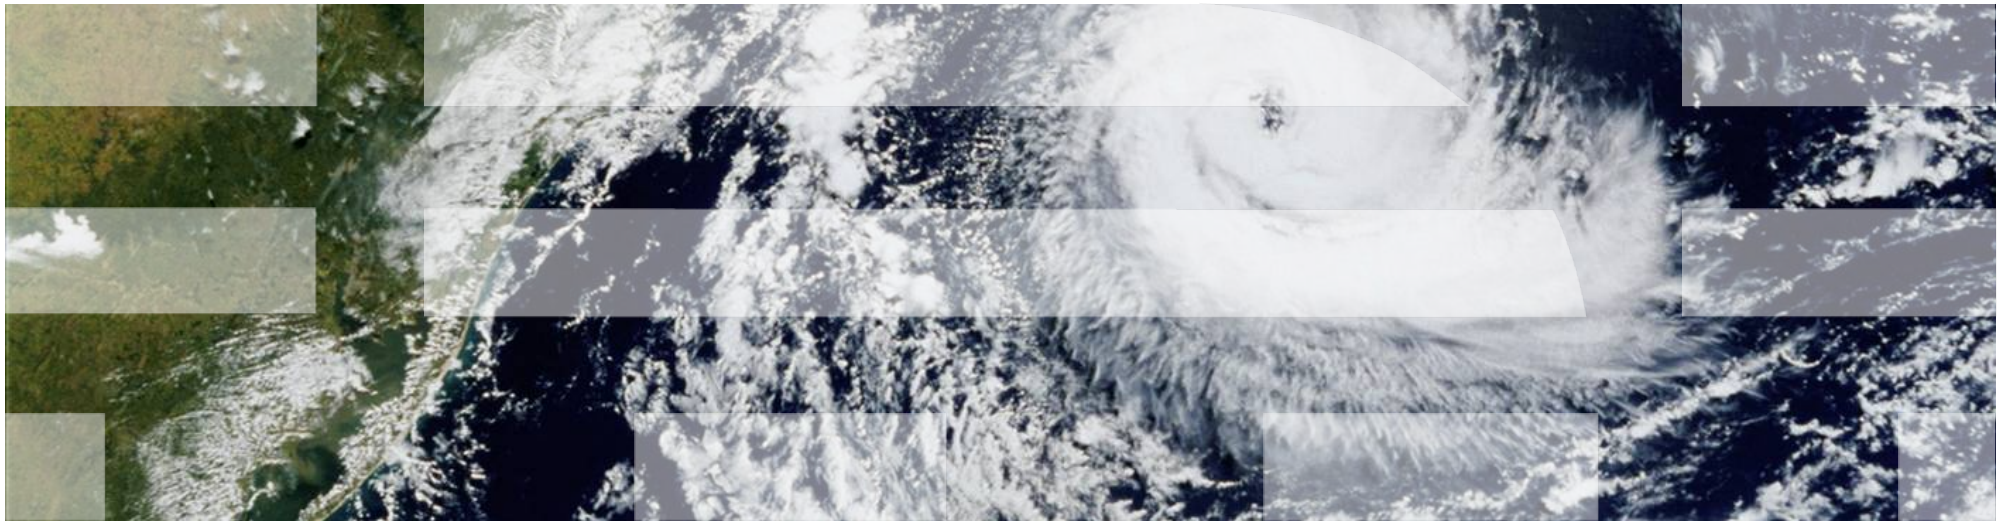


Programming Techniques for Supercomputers

(Heterogeneous and Distributed systems)

– Aditya Nitsure (IBM India)



OpenMP (Open Multi-Processing)

(<https://www.openmp.org>)

Take away

- What is OpenMP?
- How OpenMP programming model works ?
- How to write, compile and run OpenMP program?
- What are various OpenMP directives, runtime library API & Environment variables?
- How OpenMP program executes on accelerators (GPUs)?
- What are the best practices to write OpenMP program?

Content – day1 (2rd OCT)

- Introduction
- Compiling OpenMP program
- OpenMP CPU directives
 - Parallel construct
 - SIMD construct
 - Combined construct
 - Work sharing construct
 - synchronization construct
 - Tasking construct
- Hands-on – Experiment with OpenMP directives

Content – day2 (3rd OCT)

- Parallel constructs for GPU programming
- OpenMP Clauses
- Data sharing attributes
- Runtime library
- Environment variables
- Best practices
- Hands-on – Implement OpenMP directives in matrix multiplication program

Introduction (1)

- Governed by OpenMP Architecture Review Board (ARB)
- Open source, simple and up to date with the latest hardware development
- Specification for shared memory parallel programming model for Fortran and C/C++ programming languages
 - ♦ compiler directives
 - ♦ library routines
 - ♦ environment variables
- Widely accepted by community – academician & industry

Introduction (2)


- Shorter learning curve
- Supported by many compiler vendors
 - GNU, LLVM, AMD, IBM, Cray, Intel, PGI, OpenUH, ARM
 - Compilers from research labs e.g. LLNL, Barcelona Supercomputing Centre
- Easy to maintain sequential, CPU parallel and GPU parallel versions of code.
- Used in application development from variety of fields and systems
 - Aeronautics, Automotive, Pharmaceuticals, Finance
 - Supercomputers, Accelerators, Embedded multi-core systems

Specification History

- OpenMP 1.0 (1997 (Fortran) / 1998 (C/C++))
 - First release of Fortran & C/C++ specification (2 separate specifications)
- OpenMP 2.0 (2000 (Fortran) / 2002 (C/C++))
 - Addition of timing routines and various new clauses – firstprivate, lastprivate, copyprivate etc.
- OpenMP 2.5 (2005)
 - combined standard for C/C++ & Fortran
 - Introduce notion of internal control variables (ICVs) and new constructs – single, sections etc.
- OpenMP 3.0/3.1 (2008 – 2011)
 - Concept of *task* added in OpenMP execution model
- OpenMP 4.0 (2013)
 - Accelerator (GPU) support
 - SIMD construct for vectorization of loops
 - Support for FORTRAN 2003
- OpenMP 4.5 (2015)
 - Significant improvement for devices (GPU programming) & Fortran 2003
 - New taskloop construct
- **OpenMP 5.0 (2018)** – compilers are not yet support full set of features of this standard !
 - Extended memory model to distinguish different types of *flush* operations
 - Added *target-offload-var* ICV and OMP_TARGET_OFFLOAD environment variable
 - Teams construct extended to execute on host device

OpenMP Programming (1)

- The OpenMP API uses the fork-join model of parallel execution.
- The OpenMP API provides a relaxed-consistency, shared-memory model.
- An OpenMP program begins as a single thread of execution, called an initial thread. An initial thread executes sequentially.
- When any thread encounters a parallel construct, the thread creates a team of itself and zero or more additional threads and becomes the master of the new team.



OpenMP Programming (2)

- Compiler directives
 - Specified with **pragma** preprocessing
 - Starts with **#pragma omp** (For C/C++)
 - Case sensitive
 - Any directive is followed by one or more clauses
 - Syntax **#pragma omp directive-name [clause[[,] clause] ...] new-line**
- Runtime library routines
 - The library routines are external functions with “C” linkage
 - **omp.h** provides prototype definition of all library routines
- Environment variables
 - Set at the program start up
 - Specifies the settings of the Internal Control Variables (ICVs) that affect the execution of OpenMP programs
 - Always in upper case

OpenMP Hello World

helloWorld.c

```
#include <stdio.h>
#include <omp.h>

int main()
{
    int i = 0;
    int numThreads = 0;
    // call to OpenMP runtime library
    numThreads = omp_get_num_threads();

    // OpenMP directive
    #pragma omp parallel
    {
        // call to OpenMP runtime library
        int threadNum = omp_get_thread_num();
        printf("Hello World from thread %d \n", threadNum);
    }
    return 0;
}
```

```
// Setting OpenMP environment variable
export OMP_NUM_THREADS=4
```

```
//compile helloWorld program
xlc -qsmp=omp helloWorld.c -o helloWorld
```

```
// Run helloWorld program
./helloWorld
```

Output

```
[aditya@hpcsw7 openmp_tutorial]$ ./helloWorld
Hello World from thread 0
Hello World from thread 3
Hello World from thread 1
Hello World from thread 2
```

```
[aditya@hpcsw7 openmp_tutorial]$ ./helloWorld
Hello World from thread 0
Hello World from thread 1
Hello World from thread 2
Hello World from thread 3
```

Compilation of OpenMP program

XL	GNU
<ul style="list-style-type: none"> • CPU -qsmp=omp 	<ul style="list-style-type: none"> • CPU -fopenmp
<ul style="list-style-type: none"> • GPU -qsmp=omp -qoffload -qtgtarch=sm_70 (V100 GPU) 	<ul style="list-style-type: none"> • GPU -fopenmp -foffload=nvptx-none

NVIDIA CUDA linking

-lcudart -L/usr/local/cuda/lib64

OpenMP Directives


- Parallel construct
- SIMD construct
- Combined construct
- Work sharing construct
- Master and synchronization construct
- Tasking construct
- Device construct

Parallel construct

```
#pragma omp parallel [clause[ [,] clause] ... ] new-line  
{  
}
```

```
#pragma omp parallel
```

```
{  
    int threadNum = omp_get_thread_num();  
    printf("Hello World from thread %d \n", threadNum);  
}
```

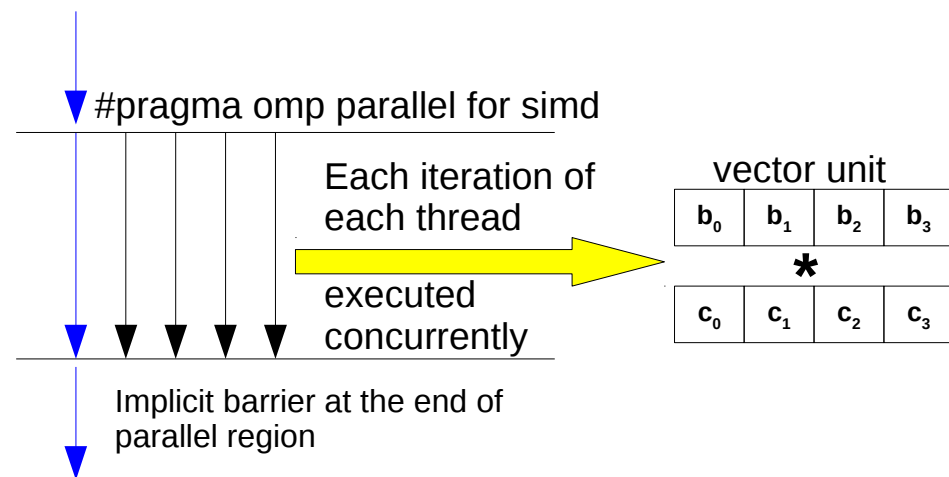


- The fundamental construct that starts parallel execution
- A team of threads is created to execute parallel region
- Original thread becomes *master* of the new team
- All threads in the team executes parallel region

SIMD construct

```
#pragma omp simd [clause[ [,] clause] ... ] new-line
{
}
```

```
#pragma omp parallel for simd
for (i=0; i<N; i++)
{
    a[i] = b[i] * c[i]
}
```



- Applied on loop to transform loop iterations to execute concurrently using SIMD instructions
- When any thread encounters a simd construct, the iterations of the loop associated with the construct may be executed concurrently using the SIMD lanes that are available to the thread.

Combined construct

- Combination of more than one construct
 - Specifies one construct immediately nested inside another construct
- Clauses from both constructs are permitted
 - With some exceptions e.g. **nowait** clause cannot be specified in **parallel for** or **parallel sections**
- Examples
 - #pragma omp parallel for
 - #pragma omp parallel for simd
 - #pragma omp parallel sections
 - #pragma omp target parallel for simd

Worksharing construct

- Distributes the execution of the associated parallel region among the members of the team
- Implicit barrier at the end
- Types
 - Loop construct
 - Sections construct
 - Single construct
 - Workshare construct (only in Fortran)

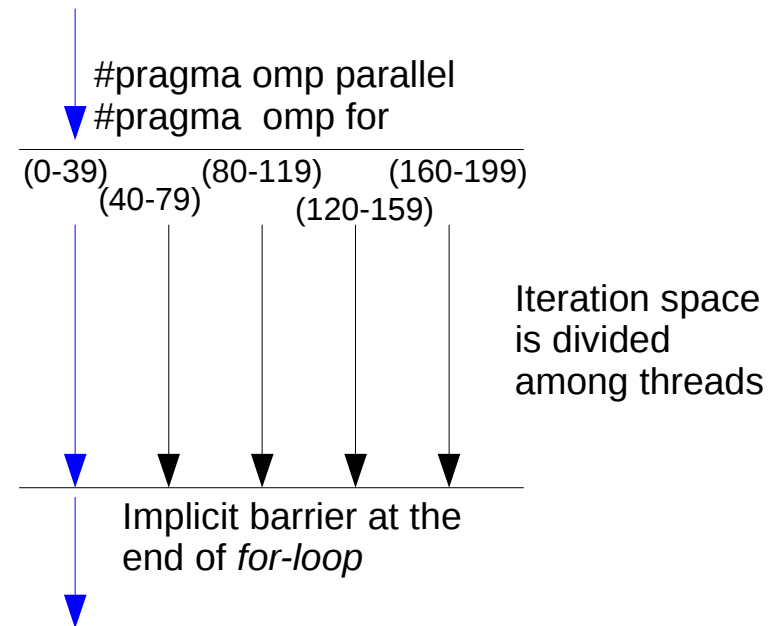
Worksharing : Loop construct

#pragma omp for [clause[[,] clause] ...] new-line

// for-loop

```
{  
}
```

```
int N = 200;  
#pragma omp parallel  
{  
    #pragma omp for  
  
    for (i=0; i<N; i++)  
    {  
        a[i] = b[i] * c[i]  
    }  
}
```



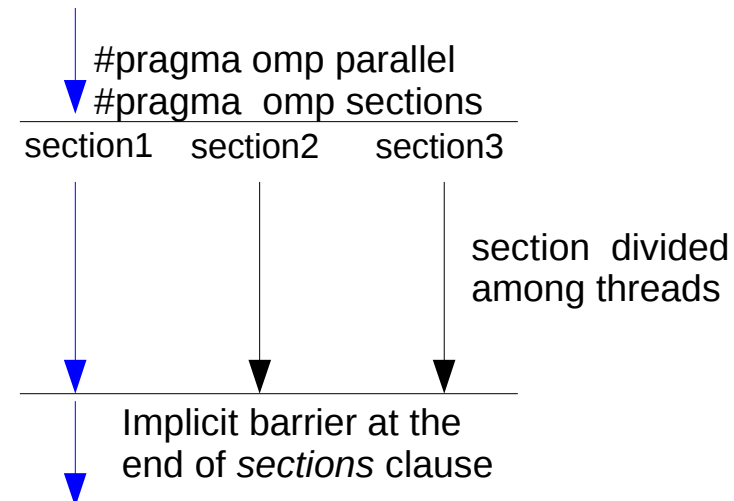
- Iterations of *for-loop* distributed among the threads
- One or more iterations executed in parallel
- Implicit barrier at the end unless *nowait*

Worksharing : Section construct

```
#pragma omp sections [clause[ [,] clause] ... ] new-line
{
    #pragma omp section new-line
    // structured block
    #pragma omp section new-line
    // structured block
}
```

- Non iterative worksharing construct
- Each section structured block is executed once
- The section structured blocks are executed in parallel (one thread per section)
- Implicit barrier at the end unless *nowait* specified

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        {
            // do_work1
        }
        #pragma omp section
        {
            // do_work2
        }
    }
}
```



Worksharing : Single construct

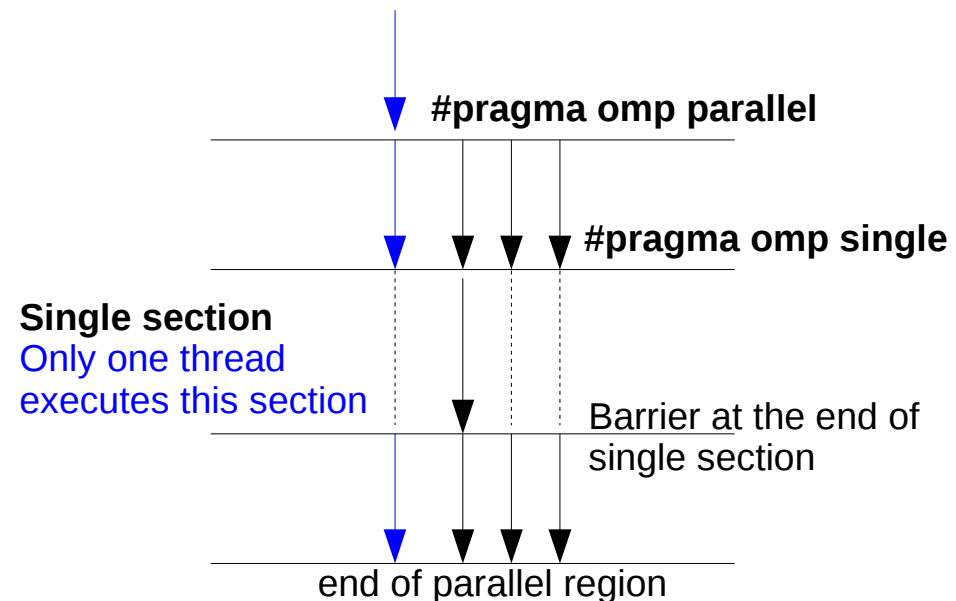
#pragma omp single [clause[[,] clause] ...] new-line

{
}

```
#pragma omp parallel
{
    // do_some_work()

    #pragma omp single
    {
        //write result to file
    }

    // do_post_work()
}
```




- Only one thread in the team executes the associated structured block
- Implicit barrier at the end
- All other threads wait until single block is finished
- *nowait* clause removes barrier at the end of single construct

Tasking construct

#pragma omp task [clause[[,] clause] ...] new-line

{
}

```
#pragma omp parallel
{
    #pragma omp single
    {
        While ( ptr != null)
        {
            #pragma omp task
            {
                //do_some_work();
            }
        } // end of while
    } // end of single
} // end of parallel
```




- Defines an explicit task
- A task is generated and corresponding data environment is created from the associated structured block
- Execution of task could be immediate or defer for later execution
- Examples : While-loop, recursion, sorting algorithms

Synchronization construct (1)

Critical construct

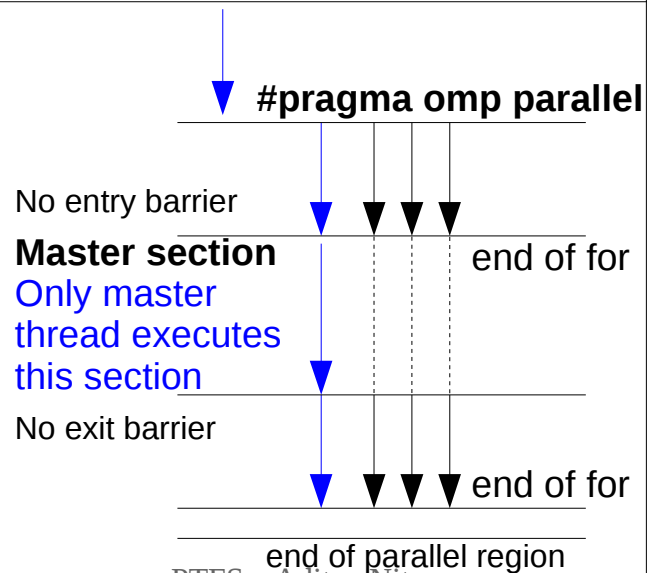
```
#pragma omp parallel
shared (sum) private (localSum)
{
    #pragma omp for
    for (i=0; i<N; i++)
    {
        localSum += a[i];
    }
    #pragma omp critical (global_sum)
    {
        sum += localSum;
    }
}
```



Master construct

```
#pragma omp parallel
{
    #pragma omp for
    // do_pre_work()

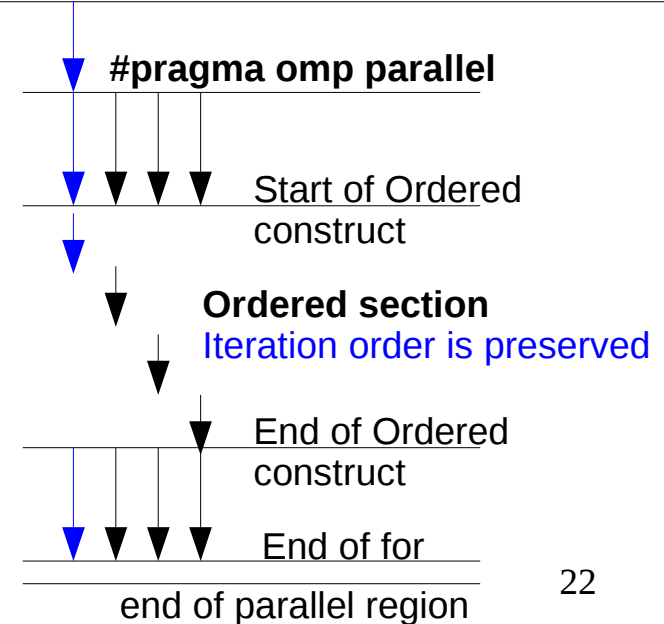
    #pragma omp master
    {
        //write result to file
    }
    #pragma omp for
    // do_post_work()
}
```



PTFS – Aditya Nitsure

Ordered construct

```
#pragma omp parallel
{
    #pragma omp for
    for (i=0; i<N; i++)
    {
        // do_some_work
        #pragma omp ordered
        {
            sum += a[i];
        }
        // do_post_work
    }
}
```



Synchronization construct (2)

- **Barrier construct**

- `#pragma omp barrier` new-line
- Stand alone directive
- “Must to execute” for all threads in parallel region
- No one continues unless all threads reach barrier

- **Taskwait construct**

- `#pragma omp taskwait` new-line
- Stand alone directive
- Waits until all child tasks completes execution before taskwait region

- **Atomic construct**

- `#pragma omp atomic` [atomic-clause] new-line
- ensures that a specific storage location is accessed atomically
- Enforces atomicity for read, write, update or capture

- **Flush construct**

- `#pragma omp flush` [(list)] new-line
- Stand alone directive
- Makes temporary view of thread’s memory consistent with main memory
- When list is specified, the flush operation applies to the items in the list. Otherwise to all thread visible data items.



Hands-on

(OpenMP directive experimentation)